

# AWS Databases

Relational, NoSQL, In-Memory, Data warehouse, Specialized

Chandra Lingam

Cloud Wave LLC

# AWS Databases



ORACLE®



Amazon DynamoDB

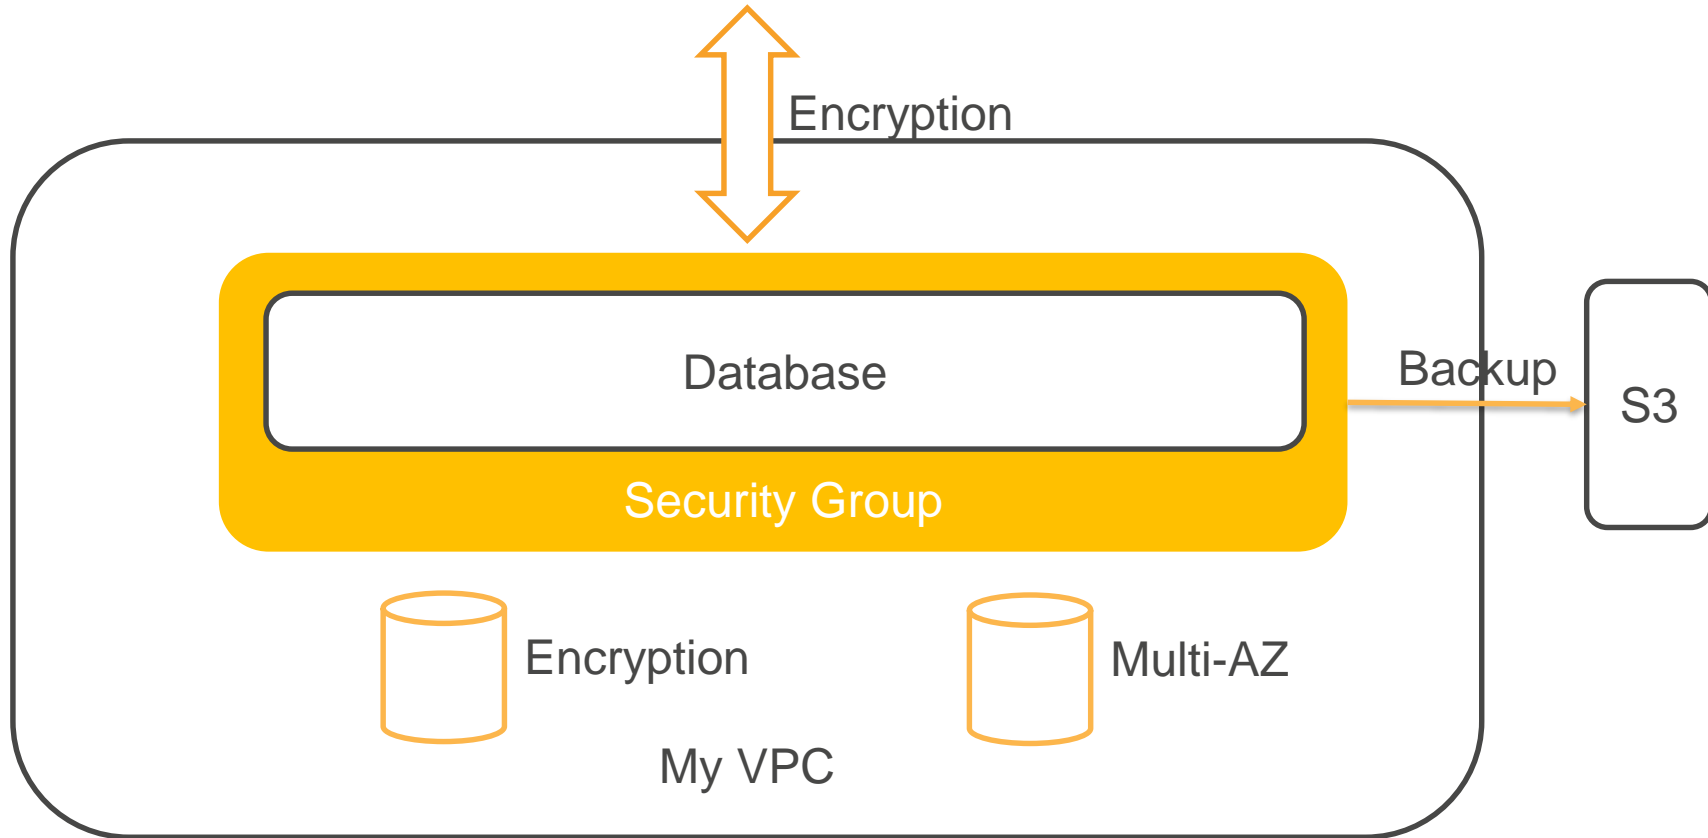


redis




Note: Not complete list

# Standard Features – AWS Databases



# DynamoDB Global Table - Multi-Region, Multi-Master



# Benefits

- Wide selection of database engines
- Fully managed
- VPC Network Isolation
- Encryption at rest using KMS
- Encryption in transit
- Automated Backup
- Highly Durable and Available – Replicated across multiple devices in Availability Zone, Region
- Multi-Region, Multi-Master (some products) – Low latency access and disaster recovery

# AWS Portfolio of Databases (1 of 2)

Service	Type of Database
RDS - Relational Database Service	Relational Database. Choice of database engines - Aurora, PostgreSQL, MySQL, MariaDB, Oracle Database, SQL Server Uses: Traditional applications, ERP, CRM, ecommerce
Redshift	Petabyte scale Data warehouse, Massively Parallel Columnar Storage, integrates with S3 data lake Uses: business intelligence, analytics, SQL to explore data lake
DynamoDB, Cassandra, DocumentDB	NoSQL Database Key-value storage, document store, consistent single digit millisecond latency at any scale Uses: high traffic web applications, ecommerce, gaming systems
ElastiCache	In-memory database - MemCached, Redis Sub-millisecond latency Uses: Caching, user session, gaming leaderboards, geospatial applications

# AWS Portfolio of Databases (2 of 2)

Service	Type of Database
Neptune	Graph Database – optimized for highly connected datasets and querying relationships Uses: Social networks, recommendation engines
Timestream	Timeseries Database – optimized for storing and querying high volume timeseries data at 1/10 <sup>th</sup> the cost of relational databases Uses: IoT applications, Industrial telemetry, DevOps
Quantum Ledger Database	Ledger Database – Blockchain based system for transparent, immutable, and cryptographically verifiable transaction log Uses: Systems of record, supply chain, banking transactions
Elasticsearch	Search database, store, analyze and correlate logs from disparate applications and systems Uses: search, infrastructure and application monitoring, Security info and event management

# Database Migration

## AWS [Database Migration Service](#) (DMS)

One-time data replication

Continuous data replication from on-premises to AWS  
([and reverse](#))

Homogeneous and Heterogeneous replication



# Summary


"The broadest selection of purpose-built databases for all your application needs"

"By picking the best database to solve a specific problem or a group of problems, you can breakaway from restrictive one-size-fits-all monolithic databases"

Reference: <https://aws.amazon.com/products/databases/>

# **Relational Database Service (RDS)**

# Relational Database



# Relational Database

General Purpose – Design a schema for any need

Rigid Schema – difficult to change

SQL – Flexible Querying System

Complex System

Scaling Challenges

# Amazon Relational Database Service (RDS)


Automates time-consuming administrative tasks (hardware, installation, patching, backup)

Production ready database in minutes

Push button scaling (CPU, Memory, Storage)

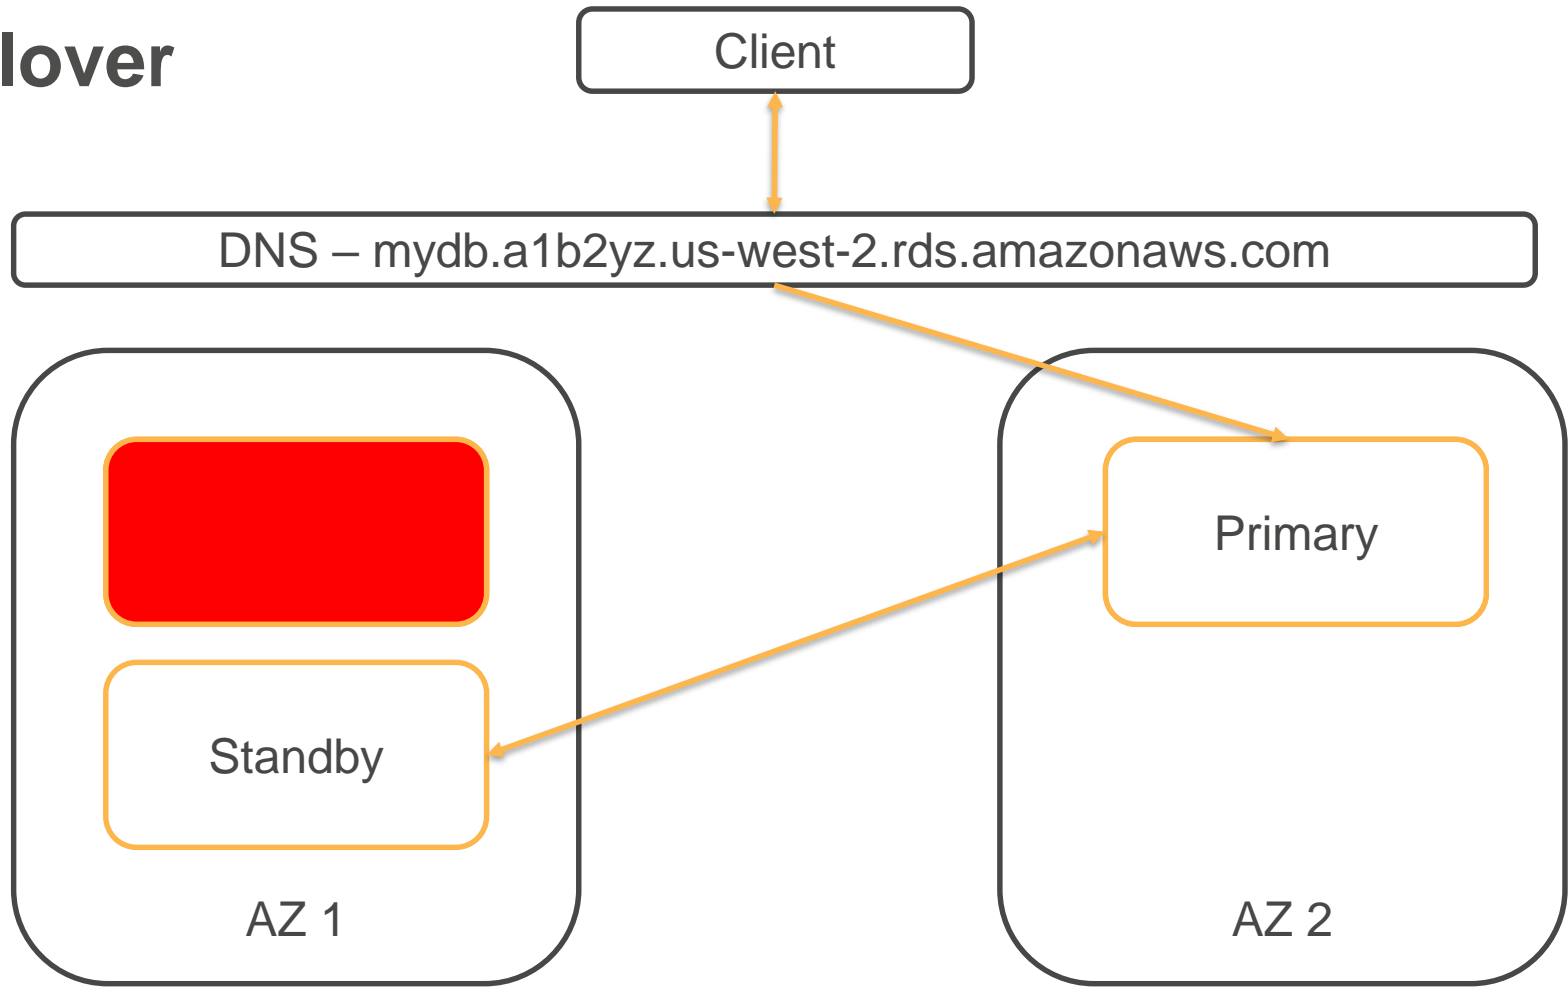
Six popular database engines: Aurora, MySQL, PostgreSQL, MariaDB, Oracle, SQL Server

# Amazon Relational Database Service (RDS)




- Connect using DNS Name
- RDS maintains mapping between DNS Name and Primary Instance
- After failover, DNS is updated to point to new primary

# Failover



# RDS - Backup and Snapshot



## Automated Backup


- Configurable for a retention up to 35 days
- Last restorable time – typically within last 5 minutes
- Point-in-time restore up to specified second (to a new instance)

## Snapshot

- User initiated
- Snapshot is kept until explicitly deleted
- Suitable for long term retention
- Copy to another region



# RDS – Read Replica



## Read Replica

- Offload Read traffic from primary instance
- Data can be stale
- One or more read replicas (depending on DB engine)

# RDS Patching

"Amazon RDS will make sure that the relational database software powering your deployment stays up-to-date with the latest patches."

You can specify a maintenance window that RDS can use for patching systems

# RDS – Scaling CPU and Memory

- Specify desired CPU and Memory configuration and RDS takes care of scaling
- Completes in a few minutes (needs to spin up new instances)
- RDS performs failover during compute scaling (interruption to client for the duration of failover)

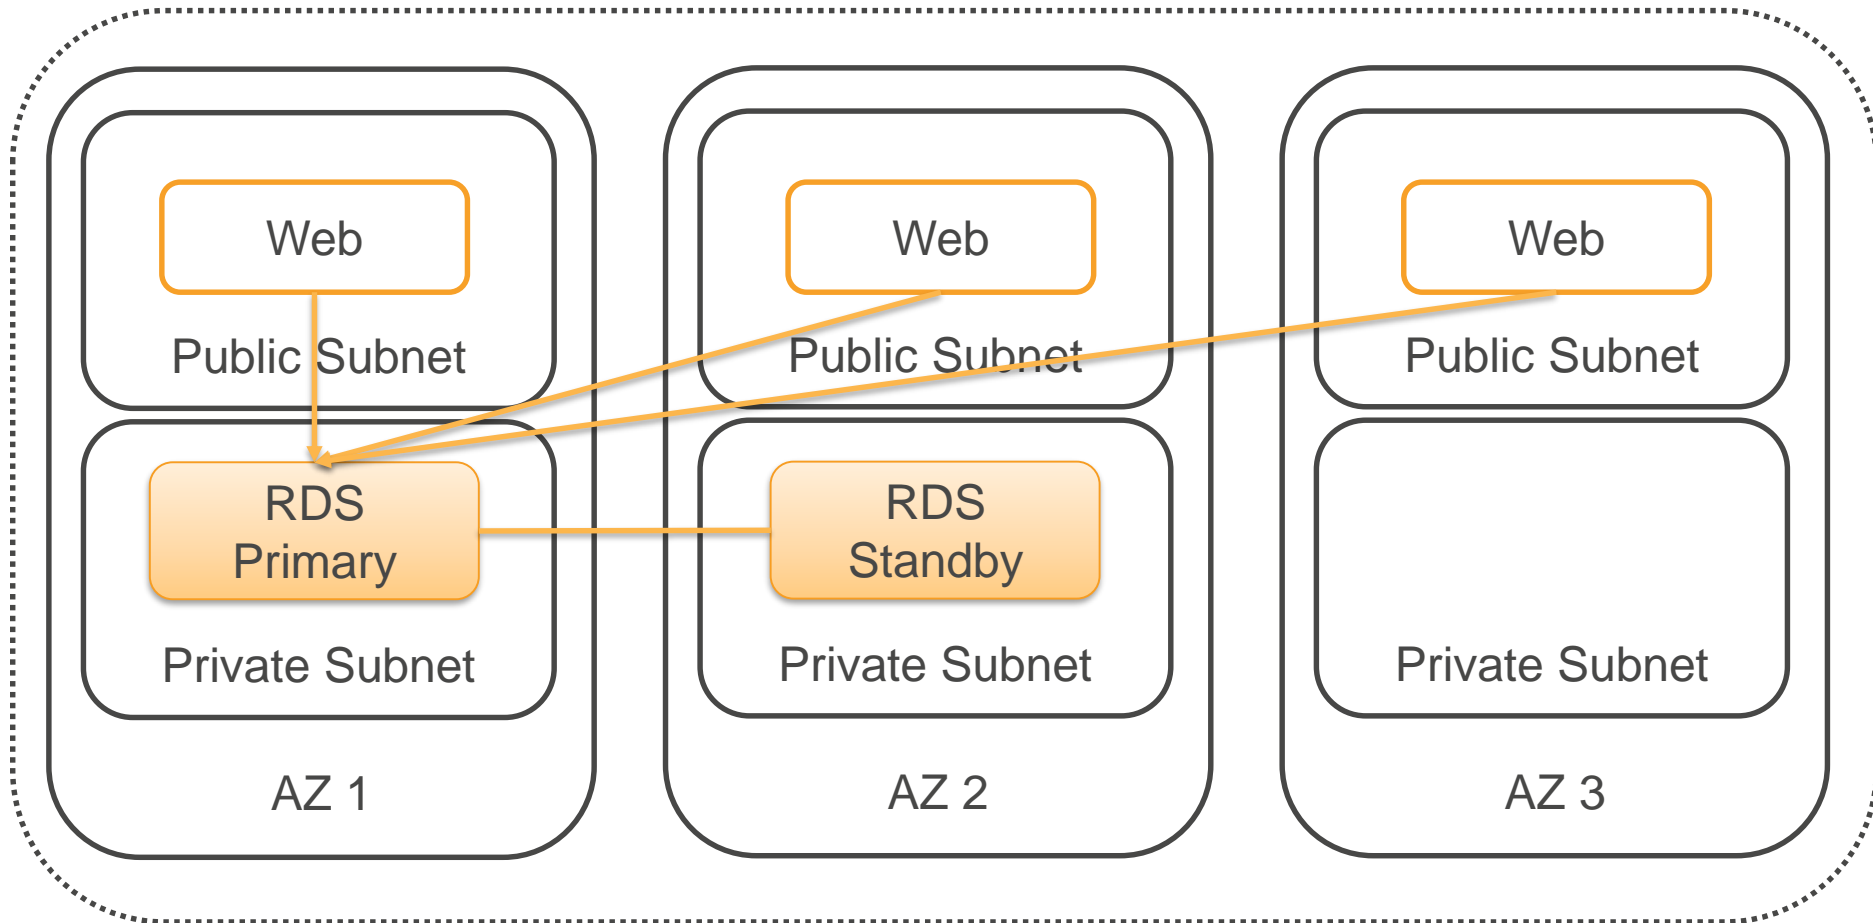
Scaling can be scheduled during next maintenance window or apply-immediate

# RDS – Storage Scaling

- Storage can be scaled without interruption (zero-downtime)
- SQL Server up to 16 TB
- Aurora up to 64 TB
- MySQL, MariaDB, PostgreSQL, Oracle up to 32 TB

Scaling can be scheduled during next maintenance window or apply-immediate

# RDS – Deployment



# RDS – Network Security

- Deploy RDS in Private Subnet (unless your requirement is a publicly accessible RDS instance)
- Configure RDS Security Group to allow access from Web Server or Application Server Security Groups
- [Assign a subnet in all Availability Zones](#) to the DB Subnet Group
  - In case of extended AZ down or some other issue, RDS may choose to launch a replacement standby instance in a different AZ
- Connect from on-premises using Amazon DirectConnect or VPN

# RDS – Permissions and Encryption

- IAM for Control Plane Access – who can create, manage, delete RDS database instances
- DB Specific User for Data Plane access – who can connect to the database, run SQL
- Optional encryption at rest using AWS Key Management Service (KMS)
- Optional encrypted connection support using SSL/TLS


# RDS – Customization, Optimization

- You can customize RDS database instance and fine tune using *DB Parameter Groups*
- RDS provides best practice guidance by analyzing configuration and usage metrics
- Use Reserved Instances for long term use (1 to 3-year terms) at substantial discount
- To prevent configuration drifts, you can use AWS Config to record and audit changes to DB instance
- For monitoring, you can use CloudWatch

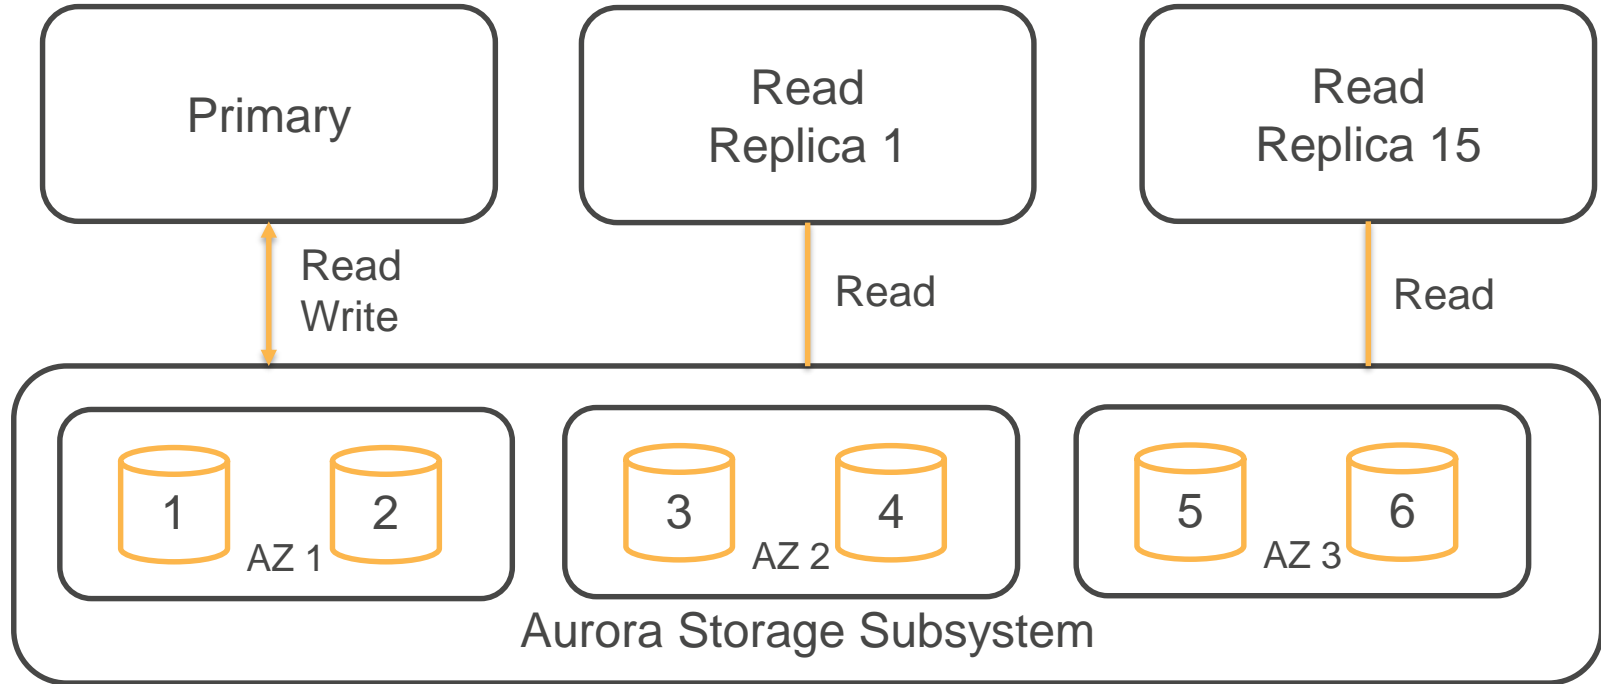


# Amazon Aurora and Aurora Serverless

# Traditional Relational Database Engine



# Amazon Aurora



# Aurora vs other Relational Databases

- Storage Subsystem that automatically maintains six copies of data across three availability zones
- Any changes made by Primary instance is replicated automatically
- Low latency Read Replica instances (lag time often in single digit millisecond)
- When the Primary fails,
  - A Read Replica is promoted as the new primary (typically under 60 seconds)
  - If Read Replica is not there, a new replacement primary is launched

# Aurora Features

- MySQL and PostgreSQL Compatibility Modes
- Up to five times faster than standard MySQL database
- Up to three times faster than standard PostgreSQL database
- Security, Availability, Reliability of commercial databases at 1/10<sup>th</sup> cost
- Support for up to 15 low latency read replicas
- [Global Database](#) - Multi-Region Replication (fast local access, disaster recovery) for globally distributed applications

# Aurora

## Cluster Endpoint

- Points to Current Primary Instance
- Suitable for Writes and Reads

mydbcluster.cluster-123456789012.us-east-1.rds.amazonaws.com:3306

## Reader Endpoint

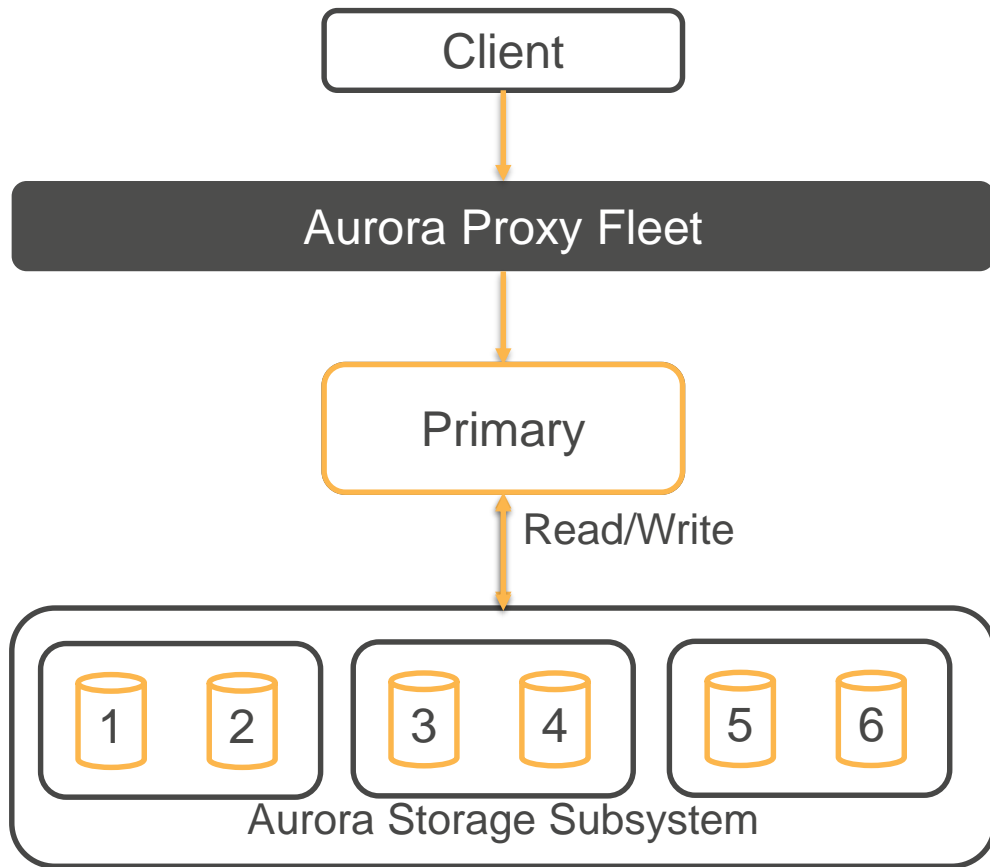
- Points to Read Replicas
- Suitable for Reads
- Multiple Read Replicas are load balanced at connection level

mydbcluster.cluster-ro-123456789012.us-east-1.rds.amazonaws.com:3306


## Instance Endpoint

- Points to Individual Aurora Instance

# Amazon Aurora Serverless



Aurora Server Warm Pool



# Aurora Serverless

- Storage and Processing are separate – scale down to zero processing and pay only for storage
- [Automatic Pause and Resume](#) – Configurable period of inactivity after which DB Cluster is Paused
  - Default is 5 minutes
  - When paused, you are charged only for Storage
  - Automatically Resumes when new database connections are requested



# Aurora Serverless

- [Aurora Serverless](#) - Suitable for use cases that are intermittent or unpredictable
- Specify Minimum, Maximum Aurora Capacity Units (ACU)
- [1 ACU](#) is ~2 GB of Memory with corresponding CPU/Network
- [Pricing](#) 1 ACU is \$0.06 per hour + Storage + I/O
- Aurora Serverless automatically scales up and down based on load
- [Scaling](#) is rapid – uses a pool of warm resources

# NoSQL Databases

DynamoDB, Cassandra, DocumentDB

# DynamoDB

- Key-value NoSQL datastore
- Flexible schema - only primary key needs to be defined
  - all columns/attributes are flexible
- Consistent performance at any scale – single digit millisecond

## Example: Movie Data

```
{  
  "year": 2013,  
  "title": "Rush",  
  "info": {  
    "directors": ["Ron Howard"],  
    "release_date": "2013-09-02T00:00:00Z",  
    "rating": 8.3,  
    "genres": ["Action", "Biography",  
               "Drama", "Sport"],  
    "actors": ["Daniel Bruhl", "Chris Hemsworth",  
               "Olivia Wilde"]  
  }  
}
```

Data Sample:


<https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GettingStarted.Python.html>

# Primary Key


Simple – Single attribute

Composite – Two attributes (partition key, sort key)

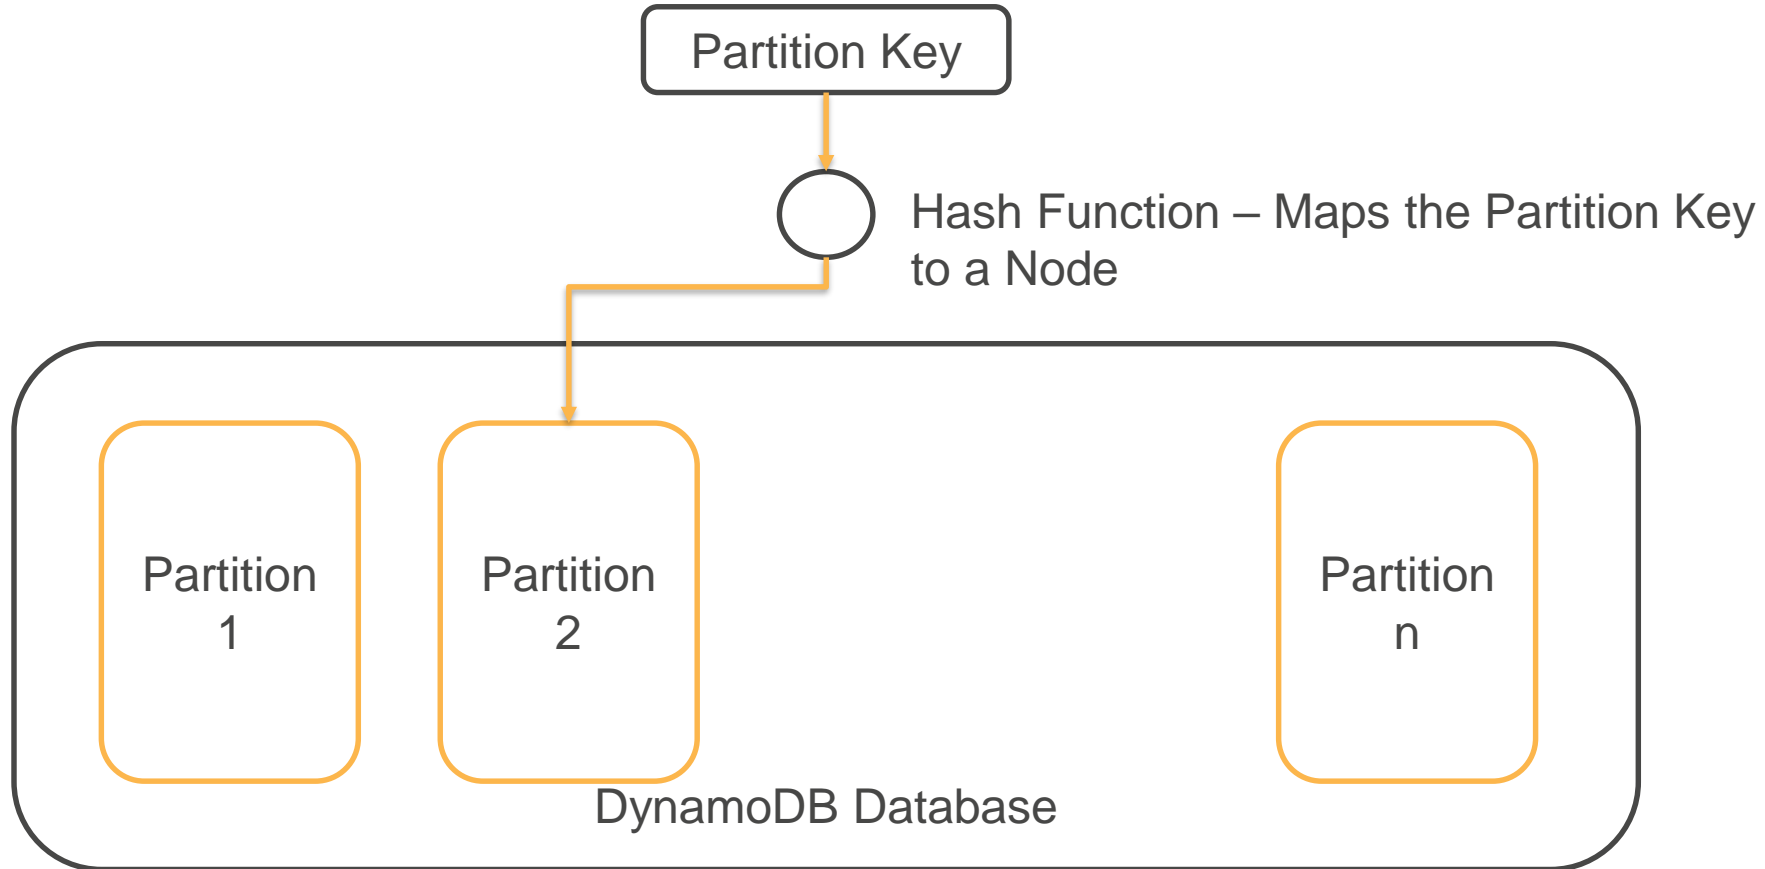
# Simple Primary Key - title



# Composite Primary Key – year, title



# Scale-Out Processing






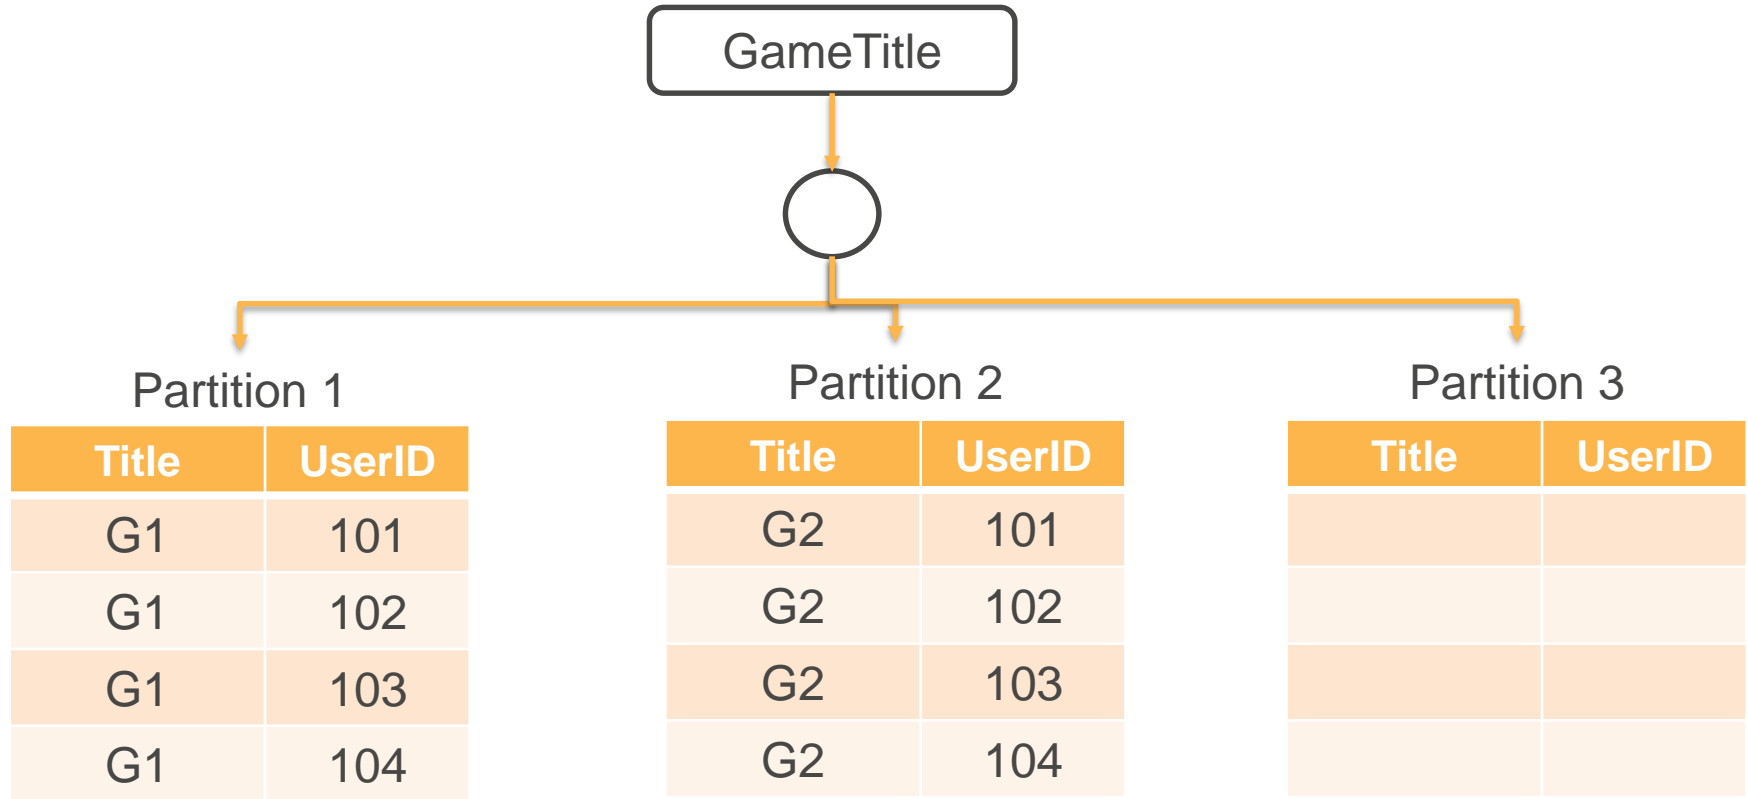
# Game Score Table

UserID	GameTitle	Country	Other attributes
101	G1	USA	
101	G2	USA	
102	G1	USA	
102	G2	USA	
103	G1	USA	
103	G2	USA	
104	G1	USA	
104	G2	USA	

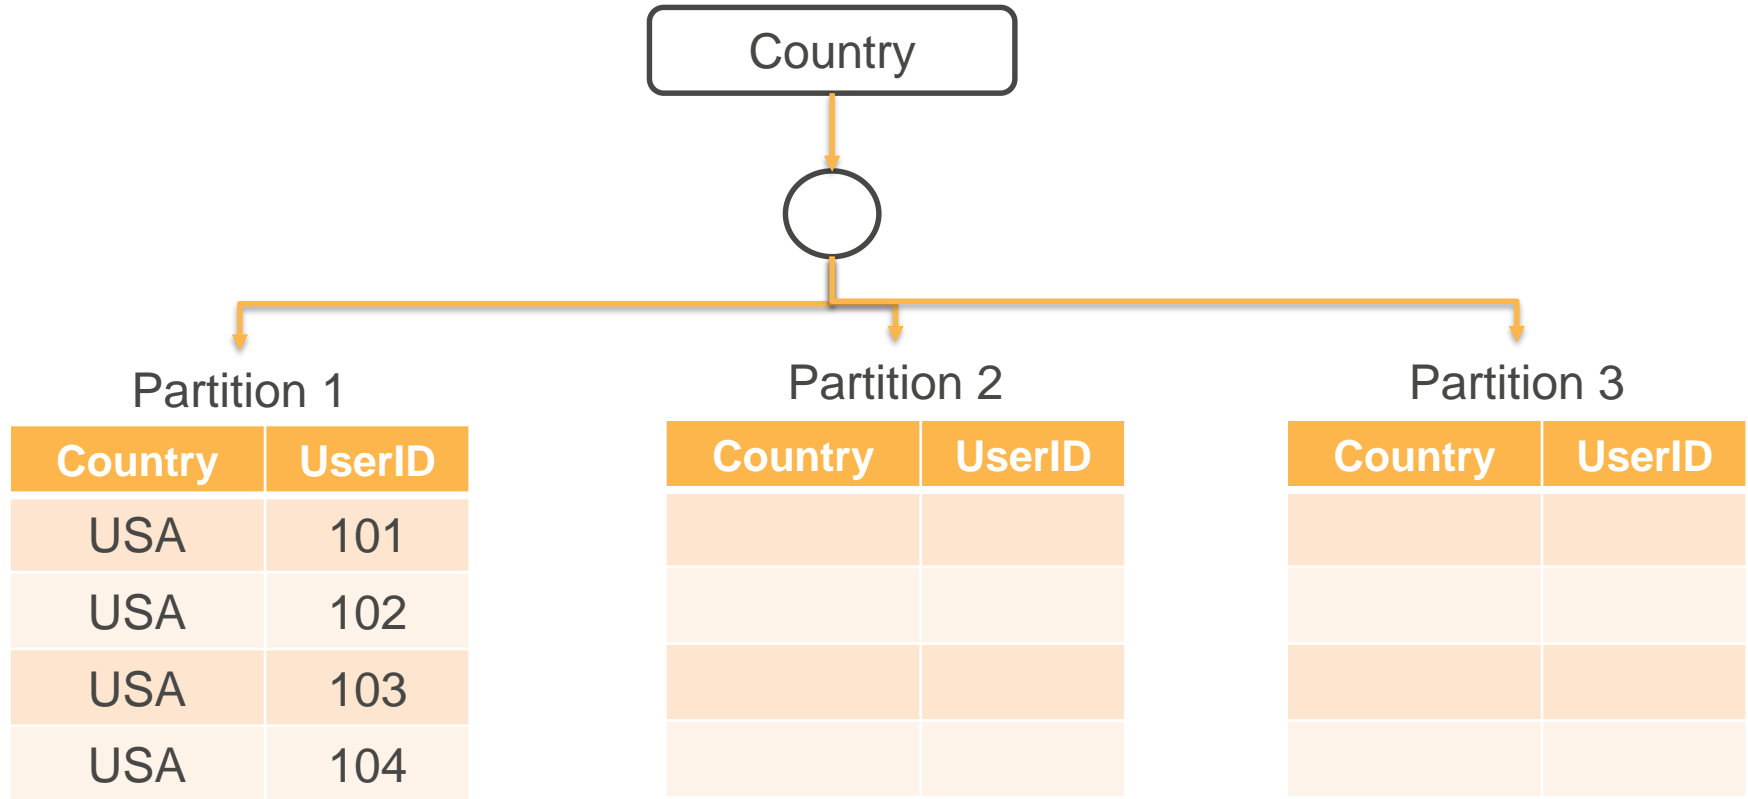
# Game Score Table – UserID, GameTitle



# Game Score Table – GameTitle, UserID




## Game Score Table – Country, UserID



# DynamoDB Features

- Automatic replication of data across multiple-availability zones in a region
- Global Tables – multi-master, multi-region replication -Fast local access across different regions
- ACID Transaction Support
- Point-in-Time Recovery – Automated Continuous Backup (35 days retention)
- On-Demand Backup/Snapshot for long term retention
- Automatic deletion of expired items – Time To Live
- Limits - Item size cannot exceed 400 KB

# DynamoDB Global Table - Multi-Region, Multi-Master



# Transactions

DynamoDB supports ACID Transactions - Atomicity, Consistency, Isolation, Durability

Transactions are useful when you want to insert, delete or update multiple items as a single logical operation

"DynamoDB provides native, server-side support for transactions, simplifying the developer experience of making coordinated, all-or-nothing changes to multiple items both within and across tables"

<https://aws.amazon.com/dynamodb/features/>

# Cassandra, DocumentDB



# Amazon Managed Cassandra

[AWS managed](#) open source Apache Cassandra

Move Cassandra workloads to AWS Cloud

Performance Benefits are comparable to DynamoDB

AWS recommends Cassandra for: industrial equipment data collection, and other use cases that require high performance and large number of columns

# Cassandra versus DynamoDB

- DynamoDB primary key is made up of single attribute partition key and optional single attribute sort key. Cassandra supports multi-column partition and sort keys
- DynamoDB max item size is 400KB – Cassandra has a [theoretical limit of 2GB](#) per column. However, general practice is not to exceed few MBs.
- Cassandra also supports large number of columns – DynamoDB even though supports large number of attributes, it is constrained by 400KB size limit per item

# Amazon DocumentDB

“Amazon DocumentDB (with MongoDB compatibility) is a fast, scalable, highly available, and fully managed document database service that supports MongoDB workloads.”

DocumentDB emulates MongoDB API and it is not true port of open source code. Currently, there is a drift in the direction of MongoDB and DocumentDB.


# ElastiCache

In-memory data store

# Amazon ElastiCache

- In-memory datastore with sub-millisecond latency
- Ideal for frequently read data, reduce read-traffic going to database, buffer high-frequency writes and periodically reconcile with backend database
  - Uses: Product reviews and rating, Caching, Session Management, Gaming leaderboards, geospatial applications
- Deploy in your VPC – Network isolation and security
- Choice of engines: Memcached, Redis

# Game Leader Boards (READs/WRITES)



# MemCached Features

- Key-value store
- Scales up to 20 nodes and 12.7 TB
- Sub milli-second latency

# Redis Features

- In-memory datastore with advanced data structures: Strings, Lists, Sorted Sets, Hash, Bit Arrays
  - Sorted Sets can be used to easily Game Leader Boards – keep a list of players sorted by rank.
- Built-in commands for [Geospatial data](#)
  - Distance between two places or persons
  - Find all places within a given distance from a point
- Sub milli-second latency
- Scales up to 250 nodes and 170 TB



# Redis High Availability Features

- Pub-Sub and Messaging  
For example, High performance chat rooms, server to server communication, social media feeds
- Read Replica across multiple Availability Zones
- Detects primary node failure and automatically promotes replica as primary
- Backup, Restore
- Export to another region
- Lua scripting support

# Amazon Redshift

# Data Warehouse - Redshift

- Peta Byte Scale Massively Parallel Relational Database
- Cluster consists of Leader Node and Multiple Compute Nodes
  - Available Storage = Storage per Compute Node X Number of Compute nodes
- [Columnar Storage](#)
- Targeted Data Compression
- Powerful SQL based Analytics
- With Redshift Spectrum - query can span tables in Redshift and files stored in S3 Data Lake